

# Genetic markers of ALS

David Lozano

December 18, 2021

## Contents

<b>1</b>	<b>Executive Summary</b>	<b>2</b>
<b>2</b>	<b>Introduction</b>	<b>2</b>
<b>3</b>	<b>Data</b>	<b>3</b>
3.1	Data sources . . . . .	3
3.2	Data cleaning . . . . .	3
3.3	Data description . . . . .	4
3.4	Data allocation . . . . .	4
3.5	Data exploration . . . . .	5
<b>4</b>	<b>Modeling</b>	<b>15</b>
4.1	Regression-based methods . . . . .	15
4.2	Tree-based methods . . . . .	21
<b>5</b>	<b>Conclusions</b>	<b>24</b>
5.1	Method comparison . . . . .	24
5.2	Takeaways . . . . .	25
5.3	Limitations . . . . .	25
5.4	Follow-ups . . . . .	26
<b>A</b>	<b>Appendix: Descriptions of features used in exploratory analysis</b>	<b>26</b>

The code to reproduce this report is available [on Github](#).

# 1 Executive Summary

**Problem.** Amyotrophic lateral sclerosis (ALS) is a fatal neurodegenerative disease for which there is no curative treatment and due to its slow onset, lacks clear markers for early diagnosis. Genetic biomarkers of ALS could revolutionize drug discovery for the disease and help clinicians with early diagnosis, thus leading to much better outcomes for patients. For these reasons I decided to investigate if gene expression levels collected from patient's blood samples could accurately classify ALS patients. Which if true, could be very useful with early diagnosis of the disease. Furthermore, I was interested to see if the genes found to be most predictive of an ALS diagnosis could give insights into the biology underlying the disease.

**Data.** My dataset was an ALS research study obtained from the National Institute of Health and submitted from the the University Medical Center Utrecht, in the Netherlands. This data set contained clinical and genetic data on over 700 patients. The explanatory variables were gene expression levels for all the genes detected in the patients blood samples. This consisted of around 29,000 genes, which I reduced to 2500 genes, essentially based on which genes were most likely to be changed between the classes. The primary response variable was patient diagnosis, which consisted of two classes, patients diagnosed with ALS and control patients.

**Analysis.** I first split my data into a training data set and a test data. I then explored the training data by assessing the characteristics of the feature variables, which were gene expression levels and the response variable, which was the diagnosis of the patient. A key aspect of this exploratory analysis involved assessing how gene expression levels differed between the ALS class and control patient class. After exploring the data, I then tuned and tested 5 different cross validated classification models: logistic regression, LASSO regression, ridge regression, elastic net regression, and random forest. Overall the lasso regression performed the test with a test missclassification error of 13.5%. The tree based, random forest model had a test missclassification error of 16%.

**Conclusions.** I found that by using solely genetic data, which was not used for the original diagnosis, penalized regressions and random forests were able accurately diagnose ALS patients vs. controls ~85% of the time. This was well above chance and suggests that genes found in easily accessible blood samples, could be important supplements to diagnosing ALS. Furthermore 2 features were present in the top 10 predictive features of the best performing models. These features were the Atp5a1 gene and AIF1 gene, which I discovered were indeed verified by scientific studies to be crucial in the onset/progression of ALS. I hope that these results provide further evidence that genetic data can be important diagnostic indicators, as well as important tools for revealing the mechanisms underlying ALS disease pathology.

## 2 Introduction

**Background.** Amyotrophic lateral sclerosis (ALS) is a fatal neurodegenerative disease that progressively deteriorates motor neurons in the brain and spinal cord. Approximately 5,000 new cases of ALS occur each year in the United States, with initial symptoms of the disease being as subtle as muscle twitches in the limbs but respiratory muscle failure commonly occurring within 2-5 years of initial symptom presentation.<sup>1</sup>. Currently there is no known cause of ALS, as it is a highly heterogeneous disease, arising from a combination of genetic and environmental factors. Due to this complexity, early diagnosis is difficult and only 2 FDA approved medications exist, which have been shown to increase survival by a few months.<sup>2</sup>.

A key development that would greatly accelerate continuing drug discovery and early diagnosis would be the discovery of a genetic biomarker for ALS. The discovery of a reliable biomarker, which could distinguish ALS patients early on, would lead to insight into which cellular and genetic pathways contribute to disease progression and thus lead to better drug targets. A blood biomarker would be the most advantageous, as it

---

<sup>1</sup>Amyotrophic Lateral Sclerosis (ALS) Fact Sheet | National Institute of Neurological Disorders and Stroke. <https://www.ninds.nih.gov/Disorders/Patient-Caregiver-Education/Fact-Sheets/Amyotrophic-Lateral-Sclerosis-ALS-Fact-Sheet/>.

<sup>2</sup>Amyotrophic Lateral Sclerosis (ALS) Fact Sheet | National Institute of Neurological Disorders and Stroke. <https://www.ninds.nih.gov/Disorders/Patient-Caregiver-Education/Fact-Sheets/Amyotrophic-Lateral-Sclerosis-ALS-Fact-Sheet/>.

is much less invasive than tissue biopsies or cerebrospinal fluid samples and would be readily accessible for the majority of patients. Past research has found reliable risk factor gene mutations in certain families with highly heritable forms of ALS, these genes seem to be involved in the mRNA processing and autophagy (waste clearance) cellular pathways.<sup>3</sup>. Unfortunately, the majority of ALS cases are not of this highly heritable form and more research is needed for finding which genes or combinations of genes could serve as reliable biomarkers in the majority of sporadic ALS cases.<sup>4</sup>.

**Analysis goals.** As this dataset is one of the largest ever collected on ALS patients, I thought it would provide a great opportunity to investigate two major questions. One question was if/how well can modern edge data science techniques reliably classify ALS patients based only on their genetic data, which is rarely taken into consideration during clinical diagnosis. The second question was, what would be the subset of genes which would prove to be the most important determinants in classifying ALS patients and what are these genes involved in.

**Significance.** Based on these goals, I hope that the results from this report provide further support that large scale genetic data can be useful in accurately diagnosing patients. Furthermore these results may be able to find novel genes/pathways which could be of importance to the disease or further confirm the importance of established ALS related genes and cellular pathways.

## 3 Data

### 3.1 Data sources

My dataset was obtained from the National Center for Biotechnology Information Gene Expression Omnibus (GEO), this is an international public repository that contains large scale genomic data submitted by individual researchers for public access.<sup>5</sup>. The submitted study I chose for this project was titled “Whole blood transcriptome analysis in amyotrophic lateral sclerosis: a biomarker study”<sup>6</sup>

This study consisted of a dataset that contained blood mRNA expression of ALS patients and non ALS patients. The dataset was collected in Utrecht, Netherlands by the University Medical Center Utrecht and was made public on April, 05, 2018. The study was designed to identify blood genetic biomarkers for ALS. The dataset contained various files of which I chose to analyze three. The first of which is the file containing the mRNA expression levels for 29,830 gene transcripts for all of the patients. These expression values were log2 transformed and quantile normalized, which is standard for the field. The second file is clinical information file for each of the patients, with information such as which patient diagnosis sex, age at diagnosis, etc. Lastly there was a file which included documentation of the gene transcript, such as the associated functions of each transcript, the exact genetic sequence, etc.. I used an R package named “GEOquery” to download these files.

### 3.2 Data cleaning

One of my main task in cleaning the data was in merging the files to get one dataset that was in the tidy format. To achieve this, I had to find a common identifier that I could use to relate the gene expression data file from each patient to their clinical information file. This identifier was provided as each patient’s geo accession id. So after I trimmed the clinical information dataset to just each patient’s clinical diagnosis and their geo accession id, I joined this tibble to the processed gene expression tibble. This resulted in my final tidy tibble that contained all the gene expression levels for each patient as a feature and the final column was each patients diagnosis status which would serve as the response variable.

---

<sup>3</sup>Balendra and Isaacs, C9orf72-Mediated ALS and FTD. <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6417666/>.

<sup>4</sup>Amyotrophic Lateral Sclerosis (ALS) Fact Sheet | National Institute of Neurological Disorders and Stroke. <https://www.ninds.nih.gov/Disorders/Patient-Caregiver-Education/Fact-Sheets/Amyotrophic-Lateral-Sclerosis-ALS-Fact-Sheet/>.

<sup>5</sup><https://www.ncbi.nlm.nih.gov/geo/info/overview.html>.

<sup>6</sup><https://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE112676>.

My other major task involved reducing the number of gene transcripts I would use for analysis, since 29,000 proved to be computationally infeasible for my resources. So I decided to use only gene transcripts that had no NA values for the associated cellular functional data. This allowed me to see which pathways and processes could be important in discriminating between the patient groups, this reduction step left me with around 12,000 gene transcripts. I then calculated the variance of each transcript in the dataset and selected the top 2,500 transcript with the most variance, which is an accepted practice in this field. and was used in page 345 of our ISLRv2 textbook <sup>7</sup>.

Lastly I cleaned the text of the data frame describing the genes and their associated function to make it presentable for reading in this format and also for wordcloud processing.

### 3.3 Data description

#### 3.3.1 Observations

My dataset had a total of 741 observations with each observation representing a patient.

#### 3.3.2 Response Variable

My response variable was patient diagnosis, which could be ALS or control. 233 of the response variables were ALS patients and 508 were control. I was mindful of a class imbalance, which I compensated for by tuning the threshold probability during the evaluation of the models.

#### 3.3.3 Features

I had 2,500 features in data set, each of which represented mRNA expression of gene transcripts, unfortunately describing these many genes in detail was unfeasible. However in the Appendix A, I provided extended features details such as the full gene name, the cellular organelle and cellular function linked to the variables, for the subset of features used in the data exploration section.

### 3.4 Data allocation

As described in the data cleaning, I decided to only choose features that had no missing values and thus I did not have to remove any observations for inputs into my model predictions. I then decided to split the dataset into training/testing data with a 80/20 split. I decided this would be agreeable with my data because each of my observations is a patient and all patients have the same gene data were selected to be analyzed. In other words it was my assumption that each observation is independent of each other, while being fairly uniform in features at the same time. I was weary of potential class imbalance as the dataset had 741 total patients, but only 233 of which were ALS patients. With the 80/20 split I ended up with ~48% of the training data being ALS patients and ~27% in the test set. I decided to keep this distribution as I thought it would be advantageous to have more ALS patients used in fitting the model than testing the model. Furthermore, I decided that I would address class imbalance, especially in my test set, during the tuning of my models. I ensured that the same training/testing datasets would be used for all the models by setting a seed for the split and saving the training and testing data as files which would then be loaded into each model.

---

<sup>7</sup>Sohil, Sohali, and Shabbir, An Introduction to Statistical Learning with Applications in R.

## 3.5 Data exploration

### 3.5.1 Response

In Figure 1, I quantified the response variable in the training data set, which was patient diagnosis and observed that 399 were ALS patients and 194 were control.

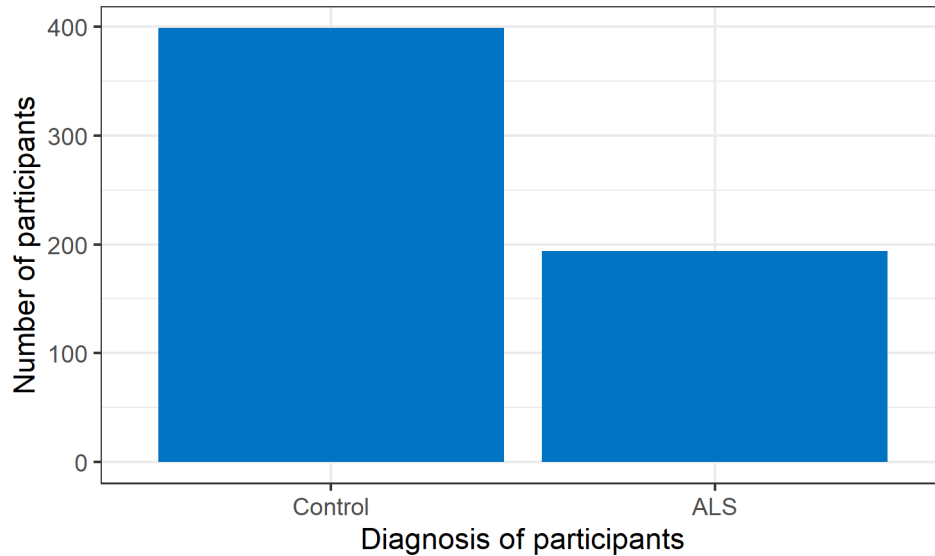


Figure 1: Distribution of ALS and non ALS patients.

### 3.5.2 Features

To understand the genetic data provided I first sought to look at the overall distribution of gene expression values in the dataset as shown in Figure 2. The expression values had been log2 transformed and quantile normalized by the original researchers, which resulted in a fairly tight range of values with most values ranging from 6 to 15. The mean expression value of this dataset was 8.6. I also looked at how the expression values for all the training features differed across the classes and could not see a noticeable difference (not shown).




Figure 2: Distribution of genetic expression values for all features in training set

Next I decided to look at the profiles a subset of genes that could be key predictors in distinguishing the two patient type. The first subset I chose to look at was a subset of the 40 features with the highest variance in the dataset. I chose to look at this subset as it is commonly thought that features with the most variance are more likely to be key in distinguishing between response classes. When I investigated how the expression values for this subset of features differed between the groups, I saw that the ALS group had slightly elevated expression levels compared to the control (Figure 3). I also was interested in seeing the correlations between the genes and if these differed across the patient groups. These results of the 5 most variable genes are shown in Figure 4 indicate that in this subset, many are correlated with each other as denoted by the asterisks. This could be due to many of the genes in this subset being involved in the same cellular pathways/functions. These correlations did not seem to be heavily altered by patient diagnosis, indicating that while some genes might be expressing at higher levels in the ALS group, their relationship with other gene may not necessarily be changing.

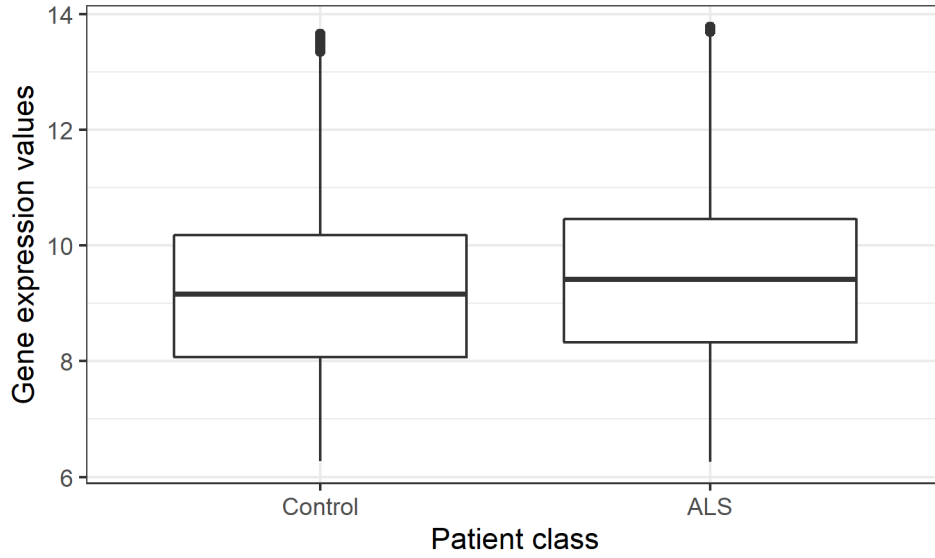


Figure 3: Difference in total gene expression between ALS and control patients

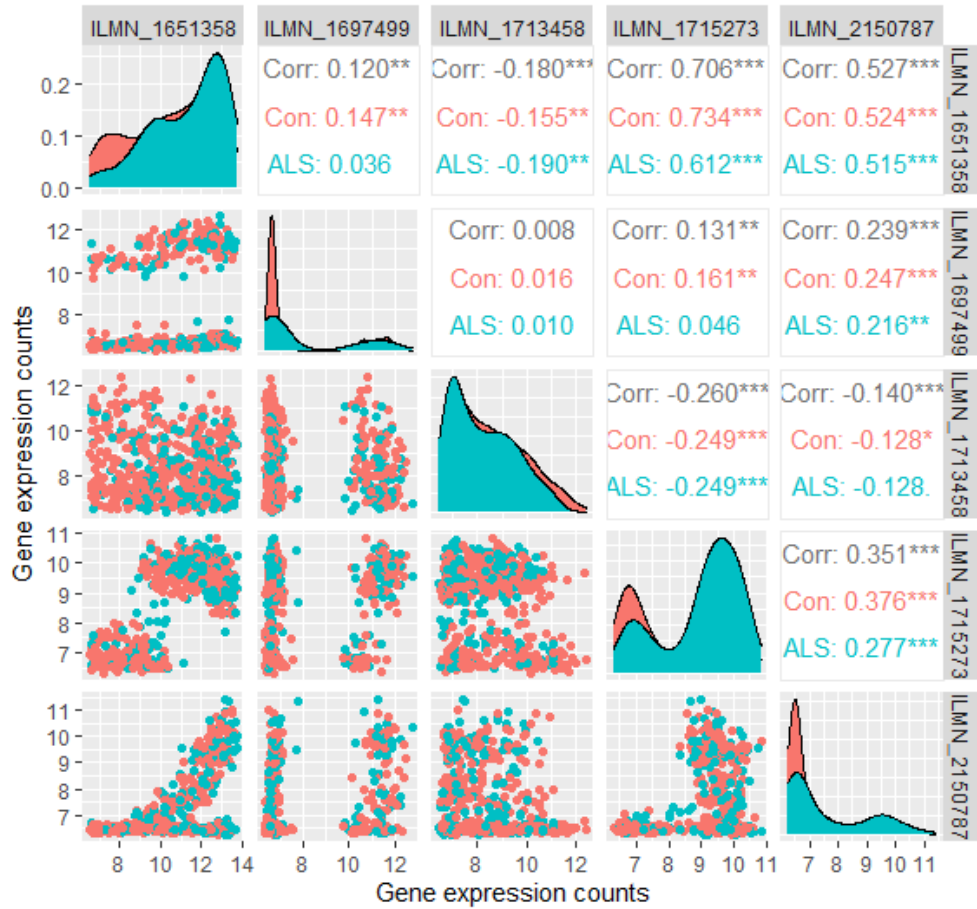


Figure 4: Correlations of top 5 most variable genes to each other and across classes

Based on these results from Figure 3, I was interested in seeing which features were the ones increased in the ALS class. Thus I calculated how the mean expression values of each feature differs between the groups. The results from Figure 5, which represents the subtracted the mean gene expression values of the ALS group from the control group, show that 32 of the 40 most variable features had increased expression levels in the ALS group.

To understand what the genes which showed increased expression in ALS group could be involved in, I decided to create visual word clouds of where in the cell (cellular organelles) these increased genes are localized to and what cellular pathways the genes are involved in (For the more details about these selected genes refer to the Appendix A). Figure 6 shows these upregulated genes were commonly localized to the cellular membrane, protein/receptor complexes and the cell's nucleus. Figure 7 shows that these upregulated genes are commonly involved in the immune and cellular defense response, as well as cellular transport. (I have also included descriptive word clouds for the downregulated genes in the Appendix A ). These associated cellular regions and functions align with many studies that have linked the motor neuron death that occurs in ALS to the inflammation response and protein recycling which heavily rely on cellular transport pathways.

8




Figure 5: Gene expression changes in ALS patients relative to controls

<sup>8</sup>Amyotrophic Lateral Sclerosis (ALS) Fact Sheet | National Institute of Neurological Disorders and Stroke. <https://www.ninds.nih.gov/Disorders/Patient-Caregiver-Education/Fact-Sheets/Amyotrophic-Lateral-Sclerosis-ALS-Fact-Sheet/>.






Figure 6: Common cell organelles for genes most upregulated ALS (from 40 most variable genes)




Figure 7: Common cellular functions for genes most upregulated ALS (from 40 most variable genes)

I found the analysis of comparing which features are increased or decreased in the ALS patients to be very informative. Thus I decided to look at the top 40 genes which were most increased/decreased in the ALS group across the entire training dataset, not just the 40 most variable genes, which is shown in Figure 8. Figure 9 revealed that the genes whose expression was most increased in the ALS patients were heavily involved in cell membrane structures, with the second most common being those involved in the cell nucleus. As for the functional information, Figure 10 revealed that these genes were also strongly involved in various regulation pathways, such as neuronal synaptic regulation, cell death regulation, and gene survival regulation. Figure 10 also showed that many of these genes were involved in protein degradation (referred to as ubiquitination) and protein transport. I found these results compelling as they were consistent with results of the 40 most variable genes and research in the field.

Next, I decided to look at the characteristics of the genes which were most decreased in the ALS class and if they differed from the upregulated genes. Figure 11 shows that these decreased ALS genes are heavily involved in the cell membrane and nucleus which was similar to the upregulated genes. However many were also being involved in the mitochondria, which was not as common in the upregulated genes. Figure 12 shows that these genes are involved in transport of various kinds, including protein transport, ion transport and cholesterol transport. This was similar to the upregulated genes, albeit not as many in the upregulated genes were involved in ion transport. Many were also involved in RNA transcription regulation and processing, which was a commonality to the upregulated genes.




Figure 8: The top 20 genes most upregulated and top 20 most downregulated in ALS patients




Figure 9: Common cell organelles for genes most upregulated in ALS (all genes)




Figure 10: Common cellular functions for genes most upregulated ALS (all genes)




Figure 11: Common cell organelles for genes most downregulated in ALS (all genes)




Figure 12: Common cellular functions for genes most upregulated ALS (all genes)

## 4 Modeling

### 4.1 Regression-based methods

#### 4.1.1 Logistic regression

I began my analysis by fitting a logistic regression on the all of the 2,500 features in the training set. I initially sought to use this fit to interpret the coefficients and extract insight into which of the genetic transcripts would be the most important predictors of patients diagnosis. As mentioned previously, the features in the data set had been log2 transformed and quantile normalized by the original researchers, as this is standard procedure for many publicly available genetic data sets. For this reason I felt it was not necessary to further transform or normalize the data.

After the logistic regression was fit, I extracted the top 10 coefficients with the most positive values as well as the the top 10 most negative coefficients. Table 1 shows these selected features, coefficients estimates,(sorted by descending value) and their associated cellular pathways for each of the coefficients. Interestingly, the variable with the highest positive coefficient was the gene NOVA1, which is localized to the cell nucleus and whose function is crucial to locomotion activity, the main deficit in ALS. Overall, many of the variables with the top positive coefficients were localized to the cell nucleus and cell membrane. Their function was commonly linked to cellular transport pathways and transcription regulation. This was in agreement with the associated functions of the upregulated genes that were explored earlier and further indicated that alterations

in cellular transport and transcription occur in ALS, which is indeed a common consensus in the field. As for the variable with the most negative coefficient values, they were also commonly localized to the plasma membrane and nucleus. Their associated functions did differ from the variables with positive coefficients in that there were more genes associated with ion transport and cell signaling. The downregulated genes previously explored were also linked to signaling and transport pathways.

Lastly I sought to address a potential class imbalance in the data set by tuning the probability threshold to see which threshold value would give the optimal missclassification error. The results shown in Table 2 demonstrate that there is no optimal threshold value, as all the values provide the same missclassification error. This result occurred because the predicted probabilities for the classes were very close to 1 or a real number multiplied by  $e^{-16}$ , meaning that essentially the probabilities already set the prediction at 0 or 1. Thus varying the thresholds would not change the result. This reveals that the logistic regression has a very poor performance in distinguishing between ALS and control subjects and is essentially guessing at chance. This poor performance is likely due to high variance overwhelming the model, as the number of features in the data set is more than triple the number of observations.

Table 1: Feature variables with the top 10 largest logistic coefficients

Estimate	Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
504	ILMN_1811363	NOVA1	neuro-oncological ventral antigen 1 (NOVA1)	nucleus	locomotory behavior	RNA binding
503	ILMN_1800430	BRF1	BRF1 homolog	transcription factor complex	regulation of transcription, DNA-dependent	RNA polymerase II transcription factor activity
419	ILMN_2123392	OR4K5	olfactory receptor	integral to membrane	G-protein coupled receptor protein signaling pathway	olfactory receptor activity
417	ILMN_1680309	ABCG8	ATP-binding cassette	membrane	transport	ATP binding
414	ILMN_2199284	ANAPC7	anaphase promoting complex subunit 7 (ANAPC7)	anaphase-promoting complex	mitosis	binding
411	ILMN_2119774	CYP2R1	cytochrome P450	membrane	electron transport	heme binding
398	ILMN_1712988	CORIN	corin	plasma membrane	morphogenesis	serine-type endopeptidase activity
385	ILMN_1761717	SLC6A1	solute carrier organic anion transporter family	membrane	transport	transporter activity
374	ILMN_2075714	ZNF284	zinc finger protein 284 (ZNF284)	intracellular	regulation of transcription, DNA-dependent	nucleic acid binding
372	ILMN_2336970	MUTYH	mutY homolog (E. coli) (MUTYH)	nucleus	metabolism	4 iron, 4 sulfur cluster binding
-401	ILMN_1811072	KCNAB1	potassium voltage-gated channel	integral to membrane	ion transport	oxidoreductase activity
-386	ILMN_1683774	IL2RA	interleukin 2 receptor	plasma membrane	immune response	interleukin-2 receptor activity
-354	ILMN_1756125	H1FNT	H1 histone family	nuclear chromatin	sperm chromatin condensation	DNA binding
-343	ILMN_2227385	SLC16A14	solute carrier family 16	membrane	transport	transporter activity
-325	ILMN_1735704	NMBR	neuromedin B receptor (NMBR)	plasma membrane	signal transduction	receptor activity
-323	ILMN_1793016	TLX2	T-cell leukemia homeobox 2 (TLX2)	nucleus	regulation of transcription, DNA-dependent	sequence-specific DNA binding
-322	ILMN_1814327	AGTR1	angiotensin II receptor	plasma membrane	circulation	angiotensin type II receptor activity
-321	ILMN_2353697	MIZF	MBD2-interacting zinc finger (MIZF)	nucleus	regulation of transcription, DNA-dependent	zinc ion binding
-317	ILMN_2134130	FLG	filaggrin (FLG)	cytoskeleton	cell differentiation	structural molecule activity
-305	ILMN_1797892	SCN11A	sodium channel	voltage-gated sodium channel complex	ion transport	sodium ion binding

Table 2: Missclassification error and related metrics for various thresholds

Threshold	Missclassification error	False positive rate	False Negative rate
0.1	0.41	0.4	0.41
0.3	0.41	0.4	0.41
0.5	0.41	0.4	0.41
0.7	0.41	0.4	0.41

#### 4.1.2 Penalized regression

Given that the logistic regression's poor performance was likely due to the high dimensional nature of the data, I decided to run penalized regressions which are more optimize for dealing with datasets with a large amount of features. Of the penalized regression models, I chose to run the ridge, LASSO (Least Absolute Shrinkage and Selection Operator), and elastic net models. All of these models went through cross-validation to select the optimal values of lambda according to the one-standard-error rule.

*Lasso:* I show in Figure 13 the cross validation plot for the Lasso, which demonstrates that the optimal model has only 46 feature. Table 3 shows the top 20 of these selected features. From the top 20 features, it can be seen that many of the selected genes are localized to nucleus and the mitochondrion, along with being important for cellular transport and RNA processing pathways.

Figure 14 essentially highlights the 10 most important features chosen by the model, and Table 3 shows the highlighted features exact coefficient values and their related functional cellular data.



I also tested various classification thresholds shown in Table 5, which shows that the standard 0.5 threshold as being the optimal one.




Figure 13: Lasso CV plot.

Table 3: Standardized coefficients for features in the lasso model based on the one-standard-error rule.

Feature ID	coefficient	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1703538	-0.42	AIF1	allograft inflammatory factor 1 (AIF1)	nucleus	response to stress	calcium ion binding
ILMN_1671661	-0.33	HSD17B7	hydroxysteroid (17-beta) dehydrogenase 7 (HSD17B7)	membrane	metabolism	oxidoreductase activity
ILMN_1796316	0.32	MMP9	matrix metalloproteinase 9 (gelatinase B)	extracellular matrix (sensu Metazoa)	peptidoglycan metabolism	zinc ion binding
ILMN_1726603	-0.32	ATP5J	ATP synthase	mitochondrion	ion transport	metal ion binding
ILMN_1684982	0.29	PDK4	pyruvate dehydrogenase kinase	mitochondrion	glucose metabolism	ATP binding
ILMN_1766010	-0.24	YARS	tyrosyl-tRNA synthetase (YARS)	extracellular space	cell motility	ligase activity
ILMN_1727184	0.18	WDR36	WD repeat domain 36 (WDR36)	small nucleolar ribonucleoprotein complex	rRNA processing	snoRNA binding
ILMN_1777220	-0.18	VCP	valosin-containing protein (VCP)	cytosol	transport	nucleotide binding
ILMN_2395375	-0.17	GABBR1	gamma-aminobutyric acid (GABA) B receptor	integral to membrane	negative regulation of adenylate cyclase activity	receptor activity
ILMN_1748770	-0.16	CKAP5	cytoskeleton associated protein 5 (CKAP5)	centrosome	RNA transport	binding
ILMN_2404112	0.16	TRDMT1	tRNA aspartic acid methyltransferase 1 (TRDMT1)	nucleus	tRNA processing	transferase activity
ILMN_1691181	0.15	TXNDC1	thioredoxin domain containing 1 (TXNDC1)	membrane fraction	cell proliferation	protein disulfide oxidoreductase activity
ILMN_1772492	-0.12	MCAAT1	mitochondrial carrier triple repeat 1 (MCAAT1)	mitochondrion	transport	binding
ILMN_1680703	0.11	MRPS15	mitochondrial ribosomal protein S15 (MRPS15)	mitochondrion	protein biosynthesis	structural constituent of ribosome
ILMN_1715963	-0.11	FBXO7	F-box protein 7 (FBXO7)	ubiquitin ligase complex	ubiquitin-dependent protein catabolism	ubiquitin-protein ligase activity
ILMN_1741143	-0.10	TXK	TXK tyrosine kinase (TXK)	cytoplasm	protein amino acid phosphorylation	ATP binding
ILMN_2313730	-0.09	RHOC	ras homolog gene family	membrane	positive regulation of I-kappaB kinase/NF-kappaB cascade	signal transducer activity
ILMN_1696485	-0.08	HNRNPAB	heterogeneous nuclear ribonucleoprotein A/B (HNRNPAB)	nucleus	mRNA editing	mRNA binding
ILMN_1745395	0.06	CTCF	COCTC-binding factor (zinc finger protein)-like (CTCF)	nucleus	regulation of transcription, DNA-dependent	zinc ion binding
ILMN_1755992	-0.05	UCP3	uncoupling protein 3 (mitochondrial)	integral to membrane	mitochondrial transport	transporter activity




Figure 14: Lasso trace plot.

Table 4: Descriptions of the 10 selected features shown in the Lasso trace plot

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1703538	AIF1	allograft inflammatory factor 1 (AIF1)	nucleus	response to stress	calcium ion binding
ILMN_1756937	STSLA4	STS alpha-N-acetyl-neuraminide alpha-2	integral to Golgi membrane	N-glycan processing	sialyltransferase activity
ILMN_1668369	CDC37	cell division cycle 37 homolog (S. cerevisiae) (CDC37)	cytoplasm	protein folding	unfolded protein binding
ILMN_1766010	YARS	tyrosyl-tRNA synthetase (YARS)	extracellular space	cell motility	ligase activity
ILMN_1661554	DIAPH1	diaphanous homolog 1 (Drosophila) (DIAPH1)	cytoskeleton	cell organization and biogenesis	actin binding
ILMN_1738632	PRKAR1A	protein kinase	cAMP-dependent protein kinase complex	regulation of transcription from RNA polymerase II promoter	cAMP-dependent protein kinase regulator activity
ILMN_1802151	OSBPPL5	oxysterol binding protein-like 5 (OSBPPL5)	cytosol	cholesterol transport	oxysterol binding
ILMN_2395375	GABBR1	gamma-aminobutyric acid (GABA) B receptor	integral to membrane	negative regulation of adenylate cyclase activity	receptor activity
ILMN_1726603	ATP5I	ATP synthase	mitochondrion	ion transport	metal ion binding
ILMN_1667430	DEGS1	degenerative spermatocyte homolog 1	membrane fraction	fatty acid desaturation	electron carrier activity

Table 5: Missclassification error and related metrics for various thresholds

Threshold	Misclassification error	False positive rate	False negative rate
0.1	0.56	0.76	0.00
0.3	0.14	0.16	0.10
0.5	0.14	0.06	0.33
0.7	0.19	0.03	0.64

*Elastic Net*: I found that the ideal elastic net model also proved to be a lasso regression, as Figure 15 shows that the cross validation plot is the exact same as Figure 13 and the optimal alpha chosen was 1. Hence the trace plot of the top selected features (Figure 15 is also identical to those chosen by the lasso. The details of the genes can be found in Table 3, and the the optimal threshold can be found in Table 5




Figure 15: Elastic net CV plot.

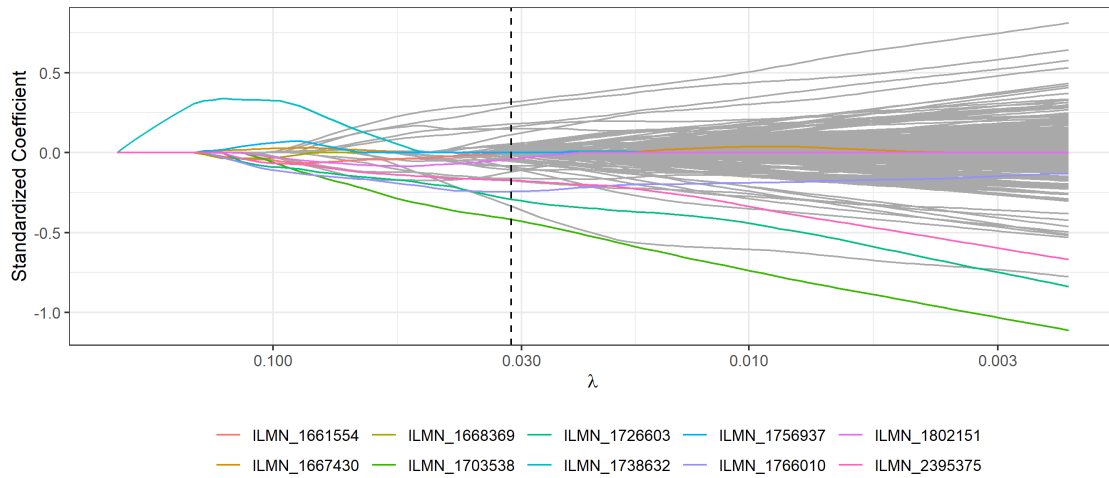


Figure 16: Elastic net trace plot.

*Ridge:* Figure 17 shows the cross validation plot for the Ridge regression. Figure 18 highlights the 10 most important features chosen by the model, and Table 6 shows the highlighted features exact coefficient values and their related functional cellular data. Similar to the lasso results, these highlighted genes are localized to nucleus and the mitochondrion, and are involved in cellular transport and RNA processing pathways. Lastly, I show the threshold tuning in Table 7, which again shows that the standard 0.5 threshold is being the optimal one.

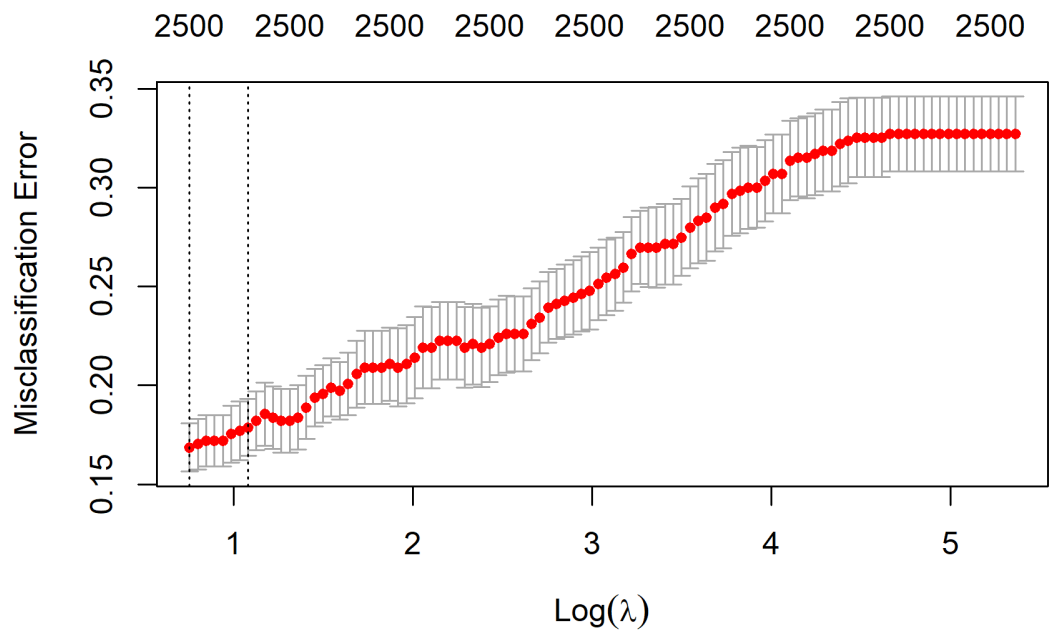


Figure 17: Lasso CV plot.

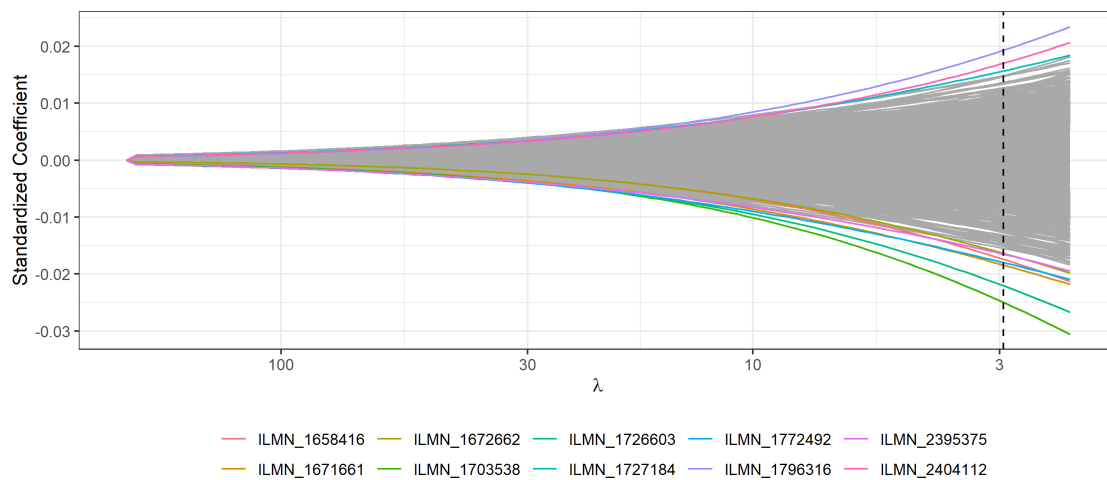


Figure 18: Ridge trace plot.

Table 6: Descriptions of the 10 selected features shown in Ridge trace plot

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1671661	HSD17B7	hydroxysteroid (17-beta) dehydrogenase 7 (HSD17B7)	membrane	metabolism	oxidoreductase activity
ILMN_1658416	MRPS18C	mitochondrial ribosomal protein S18C (MRPS18C)	mitochondrion	protein biosynthesis	structural constituent of ribosome
ILMN_1703538	AIF1	allograft inflammatory factor 1 (AIF1)	nucleus	response to stress	calcium ion binding
ILMN_2404112	TRDMT1	tRNA aspartic acid methyltransferase 1 (TRDMT1)	nucleus	tRNA processing	transferase activity
ILMN_1796316	MMP9	matrix metalloproteinase 9 (gelatinase B)	extracellular matrix (sensu Metazoa)	peptidoglycan metabolism	zinc ion binding
ILMN_1772492	MCART1	mitochondrial carrier triple repeat 1 (MCART1)	mitochondrion	transport	binding
ILMN_2395375	GABBR1	gamma-aminobutyric acid (GABA) B receptor	integral to membrane	negative regulation of adenylate cyclase activity	receptor activity
ILMN_1726603	ATP5I	ATP synthase	mitochondrion	ion transport	metal ion binding
ILMN_1727184	WDR36	WD repeat domain 36 (WDR36)	small nucleolar ribonucleoprotein complex	tRNA processing	snoRNA binding
ILMN_1672662	SLC20A1	solute carrier family 20 (phosphate transporter)	integral to membrane	transport	symporter activity

Table 7: Missclassification error and related metrics for various thresholds

Threshold	Misclassification error	False positive rate	False negative rate
0.1	0.64	0.86	0.00
0.3	0.19	0.22	0.10
0.5	0.16	0.06	0.44
0.7	0.19	0.01	0.69

## 4.2 Tree-based methods

### 4.2.1 Random forest

I chose to fit a random forest, with the aim of seeing how a model which can consider more non linear solutions performs in this complex classification task. To optimize the random forest model I tested 25 different values m values (ranging from 1-2500), which are the number of features the model considers before making a split. Figure 19 shows the results of this testing, which revealed that m=1896 provided the lowest error. Figure 20 shows how the random forest with the chosen M value performs as the number of trees increase from 1 to 500 trees, which correspond to the number of bootstraps the model aggregates over. This shows that there is a significant drop in the out of bag error after roughly 80 trees and the drop in error levels off around 400 trees. As the higher number of trees generally leads to better test performance I decided to use 500 trees for the tuned random forest model.

The features which the tuned model chose as most important to making its classification are shown in Figure 21. Table 8, shows the 15 most important variable as chosen by the mean decrease in accuracy index, which essentially is a measure of how much a model's prediction accuracy is reduced when a feature is not included in the model. These features were genes that were commonly localized to the nucleus, the mitochondria, and cell membrane. Many of these genes were involved in RNA processing as well as stress and transport pathways. Table 9, shows the 15 most important variables as chosen by the gini index, which is measure of how much the model improves its classification when it splits on that given variable. These genes had many of the same characteristics as the ones chosen by the accuracy index, with the slight difference of having more genes that were localized to the cytoplasm/extra cellular space and more involved in various protein pathways vs RNA pathways.




Figure 19: Testing m values to find optimal value

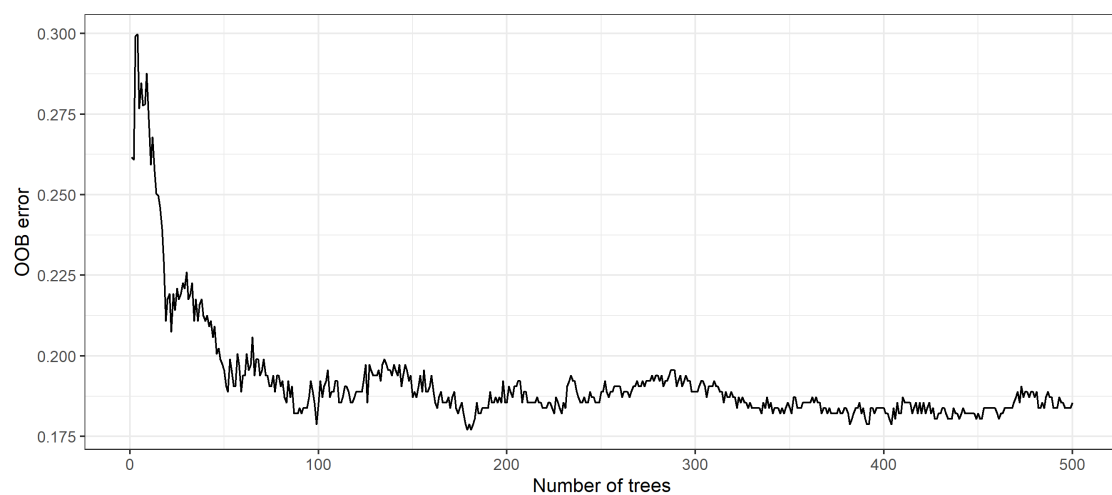


Figure 20: Optimal M value performance across 1-500 trees

## Random\_Forest\_Variable\_Importance




Figure 21: Features of most importance to random forest class predictions

Table 8: Specification of the 15 most important genes, as chosen by the mean decrease in accuracy index

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1726603	ATP5I	ATP synthase	mitochondrion	ion transport	metal ion binding
ILMN_1667430	DEGS1	degenerative spermatocyte homolog 1	membrane fraction	fatty acid desaturation	electron carrier activity
ILMN_1730611	RTN4	reticulon 4 (RTN4)	integral to membrane	negative regulation of anti-apoptosis	protein binding
ILMN_1807994	PCNP	PEST proteolytic signal containing nuclear protein (PCNP)	nucleus	protein ubiquitination	protein binding
ILMN_1703538	AIF1	allograft inflammatory factor 1 (AIF1)	nucleus	response to stress	calcium ion binding
ILMN_1756937	ST8SIA4	ST8 alpha-N-acetyl-neuraminidase alpha-2	integral to Golgi membrane	N-glycan processing	sialyltransferase activity
ILMN_2404112	TRDMT1	tRNA aspartic acid methyltransferase 1 (TRDMT1)	nucleus	tRNA processing	transferase activity
ILMN_1738632	PRKARIA	protein kinase	cAMP-dependent protein kinase complex	regulation of transcription from RNA polymerase II promoter	cAMP-dependent protein kinase regulator activity
ILMN_2307598	SLC37A3	solute carrier family 37 (glycerol-3-phosphate transporter)	membrane	transport	transporter activity
ILMN_1808148	SMCHD1	structural maintenance of chromosomes flexible hinge domain containing 1 (SMCHD1)	chromosome	chromosome organization and biogenesis	ATP binding
ILMN_1687430	EIF2B4	eukaryotic translation initiation factor 2B	eukaryotic translation initiation factor 2B complex	translational initiation	translation initiation factor activity
ILMN_1658416	MRPS18C	mitochondrial ribosomal protein S18C (MRPS18C)	mitochondrion	protein biosynthesis	structural constituent of ribosome
ILMN_1796316	MMP9	matrix metalloproteinase 9 (gelatinase B)	extracellular matrix (sensu Metazoa)	peptidoglycan metabolism	zinc ion binding
ILMN_1780291	NFAT5	nuclear factor of activated T-cells 5	nucleus	transcription from RNA polymerase II promoter	RNA polymerase II transcription factor activity
ILMN_2312340	LILRB2	leukocyte immunoglobulin-like receptor	membrane fraction	cellular defense response	receptor activity

Table 9: Specification of the 15 most important genes, as chosen by the mean decrease in Gini index

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1667430	DEGS1	degenerative spermatocyte homolog 1	membrane fraction	fatty acid desaturation	electron carrier activity
ILMN_1807994	PCNP	PEST proteolytic signal containing nuclear protein (PCNP)	nucleus	protein ubiquitination	protein binding
ILMN_1730611	RTN4	reticulon 4 (RTN4)	integral to membrane	negative regulation of anti-apoptosis	protein binding
ILMN_1756937	ST8SIA4	ST8 alpha-N-acetyl-neuraminidase alpha-2	integral to Golgi membrane	N-glycan processing	sialyltransferase activity
ILMN_2307598	SLC37A3	solute carrier family 37 (glycerol-3-phosphate transporter)	membrane	transport	transporter activity
ILMN_1738632	PRKARIA	protein kinase	cAMP-dependent protein kinase complex	regulation of transcription from RNA polymerase II promoter	cAMP-dependent protein kinase regulator activity
ILMN_1726603	ATP5I	ATP synthase	mitochondrion	ion transport	metal ion binding
ILMN_1808148	SMCHD1	structural maintenance of chromosomes flexible hinge domain containing 1 (SMCHD1)	chromosome	chromosome organization and biogenesis	ATP binding
ILMN_2404112	TRDMT1	tRNA aspartic acid methyltransferase 1 (TRDMT1)	nucleus	tRNA processing	transferase activity
ILMN_1703538	AIF1	allograft inflammatory factor 1 (AIF1)	nucleus	response to stress	calcium ion binding
ILMN_1658416	MRPS18C	mitochondrial ribosomal protein S18C (MRPS18C)	mitochondrion	protein biosynthesis	structural constituent of ribosome
ILMN_1796316	MMP9	matrix metalloproteinase 9 (gelatinase B)	extracellular matrix (sensu Metazoa)	peptidoglycan metabolism	zinc ion binding
ILMN_1687430	EIF2B4	eukaryotic translation initiation factor 2B	eukaryotic translation initiation factor 2B complex	translational initiation	translation initiation factor activity
ILMN_1666399	CDC37	cell division cycle 37 homolog (S. cerevisiae) (CDC37)	cytoplasm	protein folding	unfolded protein binding
ILMN_1766010	YARS	tyrosyl-tRNA synthetase (YARS)	extracellular space	cell motility	ligase activity

## 5 Conclusions

### 5.1 Method comparison

Table 10: Test Missclassification error of tuned models

Model	Threshold	Misclassification error	False positive rate	False negative rate
Logistic	0.5	0.405	0.404	0.410
Lasso	0.5	0.135	0.064	0.333
Elastic Net	0.5	0.135	0.064	0.333
Ridge	0.5	0.155	0.055	0.436
Random Forest	0.5	0.162	0.064	0.436

Table 10 shows the test missclassification rate at the optimum threshold, and false positive/negative rates for all the methods considered. As an model class the penalized regressions, which were the Lasso and ridge performed the best. This was slightly surprising, as I predicted the random forests to have the best performance, as I thought there would be more non linearity in the high dimensional data. The fact that the penalized regressions perform the best, indicates that reducing the dimensions/complexity of the features was most key to accurate classification in this dataset.

Thus the model that reduced the features the most, the Lasso regression, slightly outperformed the ridge by having the best overall test error of 0.135 or 13.5%. Furthermore, the lasso had the best false negative rate (33%) of all the models, indicating its algorithm reduced the amount of ALS patients which were classified as control. However the ridge had the lowest false positive rate at 5.5% (besides the logistics, which performed poorly overall), indicating that it was the best at properly distinguishing control patients from ALS patients. It seems though that it is difficult to develop an model which maintains the false positive performance of the ridge, and the false negative performance of the LASSO, since the elastic net’s attempts to do this resulted in a full lasso model. The random forest model did perform much better (missclassification rate 16.2%) than the standard logistic regression model (missclassification rate 40.2%), but it did not outcompete the penalized regressions. Lastly given that all the regression models performed best at the standard threshold of 0.5 this indicates that the class imbalance was likely not a key driver of the missclassification error.

There were general themes of which types of genes were chosen as top predictors by the models. The most common were genes localized to the cell membrane and cell nucleus. These genes were often involved in cellular transport of stress responses and RNA regulation. Slightly less common were genes in involved in the mitochondria and cytoplasm, which were commonly involved in stress responses, ion and protein transport. Surprisingly out of the thousands of genes there were two specific genes which were present in the lasso, ridge and random forest as top 15 predictor, these were the following: ATP5I, AIF1 (Table 11). Moreover, these 9 genes were present as top 15 predictors in two of the three top models: CDC37, YARS, GABBR1, DEGS1, MRPS18c, TRDMT1, MMP9, SLC20A1, PRKAR1A (Table 11).

Table 11: Genes which are top 15 predictors in multiple models

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1658416	MRPS18C	mitochondrial ribosomal protein S18C (MRPS18C)	mitochondrion	protein biosynthesis	structural constituent of ribosome
ILMN_1667430	DEGS1	degenerative spermatocyte homolog 1	membrane fraction	fatty acid desaturation	electron carrier activity
ILMN_1668369	CDC37	cell division cycle 37 homolog (S. cerevisiae) (CDC37)	cytoplasm	protein folding	unfolded protein binding
ILMN_1672662	SLC20A1	solute carrier family 20 (phosphate transporter)	integral to membrane	transport	symporter activity
ILMN_1703538	AIF1	allograft inflammatory factor 1 (AIF1)	nucleus	response to stress	calcium ion binding
ILMN_1726603	ATP5I	ATP synthase	mitochondrion	ion transport	metal ion binding
ILMN_1738632	PRKAR1A	protein kinase	cAMP-dependent protein kinase complex	regulation of transcription from RNA polymerase II promoter	cAMP-dependent protein kinase regulator activity
ILMN_1766010	YARS	tyrosyl-tRNA synthetase (YARS)	extracellular space	cell motility	ligase activity
ILMN_1796316	MMP9	matrix metalloproteinase 9 (gelatinase B)	extracellular matrix (sensu Metazoa)	peptidoglycan metabolism	zinc ion binding
ILMN_2305373	GABBR1	gamma-aminobutyric acid (GABA) B receptor	integral to membrane	negative regulation of adenylate cyclase activity	receptor activity
ILMN_2305375	GABBR1	gamma-aminobutyric acid (GABA) B receptor	integral to membrane	negative regulation of adenylate cyclase activity	receptor activity
ILMN_2404112	TRDMT1	tRNA aspartic acid methyltransferase 1 (TRDMT1)	nucleus	tRNA processing	transferase activity



## 5.2 Takeaways

My results show that analyzing genetic data can be of great value in diagnosing a complex disease such as ALS, as the best model was able to distinguish ALS patients and control patients around 87% of the time, using only genetic data. In reality, a clinical diagnosis of ALS involves multiple hospital visits, usually over a protracted period of time and is only distinguishable from other possible diseases once the physical symptoms are severe. The majority of these patients never have any kind of genetic screening to help with the clinical diagnosis. Even this data set which is one of the largest collected on ALS patients, had only collected hundreds of ALS patient over years. This is pails in comparison to the thousands of clinical ALS diagnoses that are registered every year in just the U.S.

With genetic screening becoming more accessible, this data suggests that regular genetic screenings from the blood samples from patients could serve as a very accessible and useful supplement to help clinicians make earlier diagnoses. However I think this data also suggests that diagnoses should not be made solely on genetic data, as an missclassification rate of 87% is still too low to be reliable in a clinical setting. But it could be a key supplement for expert clinicians treating the disease. Moreover, if genetic screenings did become more prevalent this would allow for larger datasets and more accurate genetic profiles for ALS. These larger data sets would not only help with diagnosing ALS but also with better understanding of the biological processes underlying disease and possible treatments.

To that end, my findings also revealed that two genes were especially key in helping the 3 best performing models classify between ALS patients and control patients. These two genes could be insights into the biology driving the disease. The first of these genes was the *Atp5a1* gene, which is important for mitochondrial function. Interestingly a 2019 scientific study in a leading journal, found that ALS onset in ALS mouse models of the disease was in part caused by abnormal *Atp5a1* expression and mitochondrial defects via this dysfunction in the pathway<sup>9</sup>. The second gene was the *AIF1* gene, which is involved in the cellular inflammatory stress response. For this gene, there were also studies in both humans and mice, which linked *AIF1* to being involved in ALS. It was especially important for the degeneration of neurons in ALS via the glia inflammatory pathway<sup>10 11</sup>. Overall the ability of the models to extract these two genes, from the thousands in the dataset, which are verified as causal contributors to ALS onset and progression is very exciting. Perhaps with larger datasets or more complex models more genes such as these or even novel genes, which are crucial to the disease, can be discovered and used to develop better treatments.

## 5.3 Limitations

### 5.3.1 Dataset limitations

A key limitation of this dataset is that the gene expression values had undergone preprocessing and normalization steps at the discretion of the original researchers. It is entirely conceivable that if a different normalization pipeline had been used on the data, that would lead to different features being selected as important by the predictive models or even ultimately changing model performance. A second limitation of that dataset is the limited number of patients. While 700 observations is well above the requirement for this report, when there are 29,000 or even 2,500 features, that leads to high dimensional data, which requires careful analysis and model tuning. The problem caused by this high dimensional data was clearly shown by the poor performance of the logistic regression model. This model was clearly overwhelmed with the variance introduced by have some many more features than observations and a dataset with less feature would likely been better for a logistic regression. Lastly there was a small class imbalance, as there were 233 ALS patients and 508 control patients in the entire dataset and this distribution was even more imbalanced

<sup>9</sup>Choi et al., "C9ORF72-ALS/FTD-Associated Poly(GR) Binds Atp5a1 and Compromises Mitochondrial Function in Vivo."

<sup>10</sup>Tam, Oliver H., Nikolay V. Rozhkov, Regina Shaw, Duyang Kim, Isabel Hubbard, Samantha Fennessey, Nadia Propp, et al. "Postmortem Cortex Samples Identify Distinct Molecular Subtypes of ALS: Retrotransposon Activation, Oxidative Stress, and Activated Glia." *Cell Reports* 29, no. 5 (October 29, 2019): 1164-1177.e5. <https://doi.org/10.1016/j.celrep.2019.09.066>.

<sup>11</sup>Peng et al., "Loss of TDP-43 in Astrocytes Leads to Motor Deficits by Triggering A1-like Reactive Phenotype and Triglycid Dysfunction."

in the test set used here. Fortunately in this case it seemed to not be a major issue, as the best threshold was still 0.5 but it is still not ideal to have uneven classes.

### 5.3.2 Analysis limitations

Due to the computational costs of analyzing all the 29,000 features, I was only able to use 2,500 of the most variable features in my analysis. Not being able to use all the gene feature data available was in my opinion was the biggest limitation in my analysis. While selecting a subset of genes based on variance is an accepted method, it is possible that I was missing features that could have led to better performance or been more predictive than the 2,500 selected features. However even with 2,500 features, I was limited in the analysis I could do because of computational costs. For example it was not possible to visualize or compute the relationships of all 2,500 to each other or to the response variable. For this reason, I could only include subsets of 40 genes in my exploratory analysis, because creating correlation plots for 2,500 was not possible for my computer. Even if it was computationally feasible these plots would not have been interpretable. Also it was limiting to not know the details of all the features in my data set, this made comparing which features were most important across models much more difficult than in datasets with less than 100 features. Overall the high dimensionality of the data proved to be very difficult analysis problem, as it meant many viable features were excluded, interpretability of features/feature importance was lost and models took much longer to run.

## 5.4 Follow-ups

As a follow up to overcome the data limitations, I would recommend adding more ALS datasets in order to increase the number of observations and even the classes. It could also be interesting to access the unprocessed files and experiment with different normalization methods of the data to see if the results are altered. As for analysis follow ups I would recommend trying models that are highly adept at high dimensional data such as neural networks of various kinds. Lastly it would be very informative to see if selecting a subset of the data by some other dimensionality reduction technique ,such as PCA, would lead to similar results.

## A Appendix: Descriptions of features used in exploratory analysis

Table 12 shows the details of the 40 most variable genes in the training data set that provide insight into the gene function and associations.

These details include:

1. The feature ID, which is the name used in most of the feature plots in the report.
2. The gene acronym, the most common way to refer to a gene in the field.
3. The extended gene name from which the acronym arises.
4. The associated cellular component which describes a cellular organelle the gene has been localized to. In reality genes are often expressed in many organelles but for simplicity/comprehension only 1 is shown for each gene
5. The associated cellular process, this aims to cover the general cellular pathway the genes have been linked to, such as DNA regulation.
6. The associated cellular function, this aims to be slightly more specific than pathways and state within the pathway what specific task the gene is important for, for example zinc ion binding.

Table 12: Specifications of 40 most variable gene transcripts in training set

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1651358	HBE1	hemoglobin	hemoglobin complex	transport	heme binding
ILMN_1697499	HLA-DREB5	major histocompatibility complex	integral to plasma membrane	immune response	MHC class II receptor activity
ILMN_1713458	HEZ	hemoglobin	hemoglobin complex	transport	heme binding
ILMN_1715273	PICALM	phosphatidylinositol binding clathrin assembly protein (PICALM)	coated pit	protein complex assembly	clathrin binding
ILMN_2150787	HLA-C	major histocompatibility complex	MHC class I protein complex	ciliary or flagellar motility	MHC class II receptor activity
ILMN_2356284	PFDN5	prefoldin subunit 5 (PFDN5)	nucleus	regulation of transcription, DNA-dependent	unfolded protein binding
ILMN_2150294	FKBP14	FK506 binding protein 14	endoplasmic reticulum	protein folding	peptidyl-prolyl cis-trans isomerase activity
ILMN_1800602	GCA	granulein	plasma membrane	membrane fusion	calcium ion binding
ILMN_1776519	RAP1GAP	RAP1 GTPase activating protein (RAP1GAP)	membrane fraction	signal transduction	GTPase activity
ILMN_1716733	MYOM2	myomesin (M-protein) 2	striated muscle thick filament	striated muscle contraction	structural constituent of muscle
ILMN_2370135	HNRNPNU	heterogeneous nuclear ribonucleoprotein U (scaffold attachment factor A) (HNRNPNU)	heterogeneous nuclear ribonucleoprotein complex	RNA processing	RNA binding
ILMN_1730454	FOLR3	folate receptor 3 (gamma) (FOLR3)	membrane fraction	folic acid transport	folic acid binding
ILMN_2363591	SDCBP	syndecan binding protein (syntenin) (SDCBP)	interleukin-5 receptor complex	synaptic transmission	protein heterodimerization activity
ILMN_2387385	IGFBP1	insulin-like growth factor binding protein 1 (IGFBP1)	extracellular space	signal transduction	insulin-like growth factor binding
ILMN_1768394	ARPC5	actin related protein 2/3 complex	Arp2/3 protein complex	cell motility	structural constituent of cytoskeleton
ILMN_1707695	IFIT1	interferon-induced protein with tetratricopeptide repeats 1 (IFIT1)	cytoplasm	immune response	binding
ILMN_1735712	KRT1	keratin 1 (epidermolytic hyperkeratosis) (KRT1)	intermediate filament	response to oxidative stress	structural constituent of cytoskeleton
ILMN_2279834	ZNF483	zinc finger protein 483 (ZNF483)	nucleus	transcription	transcription factor activity
ILMN_1796642	NCF2	neutrophil cytosolic factor 2 (65kDa)	cytosol	cellular defense response	electron carrier activity
ILMN_1684585	ACSL1	acyl-CoA synthetase long-chain family member 1 (ACSL1)	membrane	metabolism	ligase activity
ILMN_1706635	ELA2	elastase 2	extracellular region	anti-inflammatory response	serine-type endopeptidase activity
ILMN_2363586	SDCBP	syndecan binding protein (syntenin) (SDCBP)	interleukin-5 receptor complex	synaptic transmission	protein heterodimerization activity
ILMN_1766165	SNCA	synuclein	cytoplasm	anti-apoptosis	protein binding
ILMN_1669273	PPT1	palmitoyl-protein thioesterase 1 (ceroid-lipofuscinosis)	lysosome	visual perception	hydrolase activity
ILMN_2198842	KIR2DL2	killer cell immunoglobulin-like receptor	membrane	cellular defense response	receptor activity
ILMN_1671686	EPH49	erythrocyte membrane protein band 4.9 (denatin) (EPH49)	actin cytoskeleton	harbed-end actin filament capping	actin binding
ILMN_2115949	ZNF394	zinc finger protein 394 (ZNF394)	nucleus	regulation of transcription, DNA-dependent	zinc ion binding
ILMN_1797895	FFAR2	free fatty acid receptor 2 (FFAR2)	integral to plasma membrane	G-protein coupled receptor protein signaling pathway	receptor activity
ILMN_1684255	MYL4	myosin	striated muscle contraction	striated muscle contraction	structural constituent of muscle
ILMN_2380946	EIF4G2	eukaryotic translation initiation factor 4 gamma	myosin	positive regulation of translation	RNA binding
ILMN_1752347	DEFB4	defensin	extracellular region	xenobiotic metabolism	calcium channel regulator activity
ILMN_2337241	RPS15A	ribosomal protein S15a (RPS15A)	ribosome	protein biosynthesis	structural constituent of ribosome
ILMN_2173004	RAB8B	RAB8B	membrane	protein transport	GTPase activity
ILMN_1754195	RPL31	ribosomal protein L31 (RPL31)	ribosome	protein biosynthesis	structural constituent of ribosome
ILMN_1693192	PI3	peptidase inhibitor 3	extracellular region	copulation	serine-type endopeptidase inhibitor activity
ILMN_2176882	ZNF69	zinc finger protein 69 (ZNF69)	nucleus	transcription	transcription factor activity
ILMN_1710937	IFI16	interferon	nucleus	cell proliferation	binding
ILMN_1796537	FYB	FYN binding protein (FYB-120/130) (FYB)	nucleus	protein kinase cascade	receptor binding
ILMN_2347798	IFI6	interferon	membrane	anti-apoptosis	protein binding
ILMN_2155998	PSMD6	proteasome (prosome)	proteasome complex (sensu Eukaryota)	ATP-dependent proteolysis	binding

## 20 most upregulated gene transcripts in ALS patients:

Table 13 shows the details of the 20 genes which were most upregulated in ALS patients.

Table 13: Specifications of 20 most upregulated genes in ALS patients in the training set

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1758632	PRKAR1A	protein kinase	cAMP-dependent protein kinase complex	regulation of transcription from RNA polymerase II promoter	cAMP-dependent protein kinase regulator activity
ILMN_1800602	GCA	granulein	plasma membrane	membrane fusion	calcium ion binding
ILMN_1651358	HBE1	hemoglobin	hemoglobin complex	transport	heme binding
ILMN_1807994	PCNP	PEST proteolytic signal containing nuclear protein (PCNP)	nucleus	protein ubiquitination	heme binding
ILMN_1679727	CLK1	CLK-like kinase 1 (CLK1)	nucleus	cell proliferation	ATP binding
ILMN_1768197	ROD1	ROD1 regulator of differentiation 1 (S. pombe) (ROD1)	nucleus	morphogenesis	RNA binding
ILMN_2363586	SDCBP	syndecan binding protein (syntenin) (SDCBP)	interleukin-5 receptor complex	synaptic transmission	protein heterodimerization activity
ILMN_1800540	CD55	CD55 molecule	integral to plasma membrane	innate immune response	GPI anchor binding
ILMN_2336982	NPTX	neuropilin (NPTX)	integral to membrane	positive regulation of long-term neuronal synaptic plasticity	cell adhesion molecule binding
ILMN_1691181	TXNDC1	thioredoxin domain containing 1 (TXNDC1)	membrane fraction	cell proliferation	protein disulfide oxidoreductase activity
ILMN_1768394	ARPC5	actin related protein 2/3 complex	Arp2/3 protein complex	cell motility	structural constituent of cytoskeleton
ILMN_1758864	VNN2	vanin 2 (VNN2)	membrane	nitrogen compound metabolism	hydrolase activity, acting on carbon-nitrogen (but not peptide) bonds, in linear amides
ILMN_1760675	SEB1	sealing factor 3b	nucleus	nuclear mRNA splicing, via spliceosome	binding
ILMN_1771333	CD47	CD47 molecule (CD47)	plasma membrane	cell-matrix adhesion	protein binding
ILMN_1720771	STX11	syntaxin 11 (STX11)	membrane	intracellular protein transport	SNAP receptor activity
ILMN_1688526	ARL5A	ADP-ribosylation factor-like 5A (ARL5A)	intracellular	rRNA processing	GTP binding
ILMN_1718961	BNIP3L	BCL2/adrenovirus E1B 19kDa interacting protein 3-like (BNIP3L)	integral to membrane	negative regulation of survival gene product activity	protein homodimerization activity
ILMN_2344850	VPS26A	vacuolar protein sorting 26 homolog A (S. pombe) (VPS26A)	intracellular protein transport	intracellular protein transport	protein transporter activity
ILMN_2144573	CTSB	chitinase	lysosome	metabolism	chitinase activity
ILMN_1730611	RTN4	reticulon 4 (RTN4)	integral to membrane	negative regulation of anti-apoptosis	protein binding

## 20 most downregulated gene transcripts in ALS patients:

Table 14 shows the details of the 20 genes which were most upregulated in ALS patients.

Table 14: Specifications of 20 most downregulated genes in ALS patients in the training set

Feature ID	Gene Acronym	Extended Gene Name	Associated Cellular Component	Associated Cellular Process	Associated Cellular Function
ILMN_1726603	ATP5d	ATP synthase	mitochondrion	ion transport	metal ion binding
ILMN_1746290	GP9	glycoprotein IX (platelet) (GP9)	plasma membrane	blood coagulation	transferase activity
ILMN_1792538	CD7	CD7 molecule (CD7)	plasma membrane	cellular defense response	receptor activity
ILMN_1777794	PRKCSH	protein kinase C substrate 80K-H (PRKCSH)	endoplasmic reticulum	protein kinase cascade	calcium ion binding
ILMN_1681016	SPNS1	spinstar homolog 1 (Drosophila) (SPNS1)	membrane	transport	transporter activity
ILMN_1664028	CENPB	centromere protein B	chromosome	regulation of transcription	satellite DNA binding
ILMN_1755114	EIF2AK4	eukaryotic translation initiation factor 2 alpha kinase 4 (EIF2AK4)	cytosolic ribosome (sensu Eukaryota)	regulation of translational initiation	transferase activity
ILMN_1772492	MCAI1	mitochondrial carrier triple repeat 1 (MCAI1)	mitochondrion	transport	binding
ILMN_1754738	APM1	adaptor-related protein complex 1	clathrin vesicle coat	intracellular protein transport	protein binding
ILMN_1657893	TXNRD2	thioredoxin reductase 2 (TXNRD2)	mitochondrion	electron transport	oxidoreductase activity, acting on NADH or NADPH, disulfide as acceptor
ILMN_1671661	HSD17B7	hydroxysteroid (17-beta) dehydrogenase 7 (HSD17B7)	membrane	metabolism	oxidoreductase activity
ILMN_1666597	PI4KB	phosphatidylinositol 4-kinase	endoplasmic reticulum	phosphatidylinositol biosynthesis	kinase activity
ILMN_1792679	ITGA5	integrin	integral to membrane	cell adhesion	receptor activity
ILMN_1682930	SIPA1	signal-induced proliferation-associated gene 1 (SIPA1)	nucleus	negative regulation of cell adhesion	Rap GTPase activator activity
ILMN_1815500	ITPR3	inositol 1	plasma membrane	ion transport	calcium ion transporter activity
ILMN_1738239	RBM6	RNA binding motif protein 6 (RBM6)	nucleus	RNA processing	zinc ion binding
ILMN_1715968	MLL4	myeloid/lymphoid or mixed-lineage leukemia 4 (MLL4)	nucleus	transcription	transcription factor activity
ILMN_1700538	AIF1	apoptosis-inducing factor 1 (AIF1)	nucleus	response to stress	calcium ion binding
ILMN_2044471	NCR3	natural cytotoxicity triggering receptor 3 (NCR3)	integral to plasma membrane	inflammatory response	receptor activity
ILMN_1802151	OSBPL5	oxysterol binding protein-like 5 (OSBPL5)	cytosol	cholesterol transport	oxysterol binding

## Word clouds for most downregulated genes within the 40 most variable genes



Figure 22: Common cellular organelles for genes most downregulated ALS (from 40 most variable genes)




Figure 23: Common cellular functions for genes most downregulated ALS (from 40 most variable genes)